

Name: Garcia, John Carlos M.

Section: CPE22S3

Date Performed: 04/28/2024

Date Submitted: 04/28/2024

✓ Logistic Regression Analysis

Task: Predict Cervic Cancer

✓ Setup

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split
6 from sklearn.preprocessing import MinMaxScaler
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.metrics import accuracy_score
9 from sklearn.metrics import confusion_matrix
10
11 %matplotlib inline
```

```
1 pip install ucimlrepo
```

```
Collecting ucimlrepo
  Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.6
```

```
1 from ucimlrepo import fetch_ucirepo
2
3 # fetch dataset
4 cervical_cancer_risk_factors = fetch_ucirepo(id=383)
5
6 # data (as pandas dataframes)
7 X = cervical_cancer_risk_factors.data.features
8 y = cervical_cancer_risk_factors.data.targets
9
10 # metadata
11 print(cervical_cancer_risk_factors.metadata)
12
13 # variable information
14 print(cervical_cancer_risk_factors.variables)
15
```

```
{'uci_id': 383, 'name': 'Cervical Cancer (Risk Factors)', 'repository_url': 'https://archive.ics.uci.edu/dataset/383/cervical+cancer+'}
```

	name	role	type	demographic \
0	Age	Feature	Integer	Age
1	Number of sexual partners	Feature	Continuous	Other
2	First sexual intercourse	Feature	Continuous	None
3	Num of pregnancies	Feature	Continuous	None
4	Smokes	Feature	Continuous	None
5	Smokes (years)	Feature	Continuous	None
6	Smokes (packs/year)	Feature	Continuous	None
7	Hormonal Contraceptives	Feature	Continuous	None
8	Hormonal Contraceptives (years)	Feature	Continuous	None
9	IUD	Feature	Continuous	None
10	IUD (years)	Feature	Continuous	None
11	STDs	Feature	Continuous	None
12	STDs (number)	Feature	Continuous	None
13	STDs:condylomatosi	Feature	Continuous	None
14	STDs:cervical condylomatosi	Feature	Continuous	None
15	STDs:vaginal condylomatosi	Feature	Continuous	None
16	STDs:vulvo-perineal condylomatosi	Feature	Continuous	None

17	STDs:syphilis	Feature	Continuous	None
18	STDs:pelvic inflammatory disease	Feature	Continuous	None
19	STDs:genital herpes	Feature	Continuous	None
20	STDs:molluscum contagiosum	Feature	Continuous	None
21	STDs:AIDS	Feature	Continuous	None
22	STDs:HIV	Feature	Continuous	None
23	STDs:Hepatitis B	Feature	Continuous	None
24	STDs:HPV	Feature	Continuous	None
25	STDs: Number of diagnosis	Feature	Integer	None
26	STDs: Time since first diagnosis	Feature	Continuous	None
27	STDs: Time since last diagnosis	Feature	Continuous	None
28	Dx:Cancer	Feature	Integer	None
29	Dx:CIN	Feature	Integer	None
30	Dx:HPV	Feature	Integer	None
31	Dx	Feature	Integer	None
32	Hinselmann	Feature	Integer	None
33	Schiller	Feature	Integer	None
34	Citology	Feature	Integer	None
35	Biopsy	Feature	Integer	None

	description	units	missing_values
0	None	None	no
1	None	None	yes
2	None	None	yes
3	None	None	yes
4	None	None	yes
5	None	None	yes
6	None	None	yes
7	None	None	yes
8	None	None	yes
9	None	None	yes
10	None	None	yes
11	None	None	yes
12	None	None	yes
13	None	None	yes
14	None	None	yes
15	None	None	yes
16	None	None	yes

Data Frame

```
1 logistic_df = pd.concat([X, y], axis=1) #Combine both dataframes into one for more efficient manipulation of data
```

Exploration

```
1 logistic_df
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormo Contracepti
0	18	4.0	15.0	1.0	0.0	0.0	0.0	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	
2	34	1.0	NaN	1.0	0.0	0.0	0.0	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	
...	
853	34	3.0	18.0	0.0	0.0	0.0	0.0	
854	32	2.0	19.0	1.0	0.0	0.0	0.0	
855	25	2.0	17.0	0.0	0.0	0.0	0.0	
856	33	2.0	24.0	2.0	0.0	0.0	0.0	
857	29	2.0	20.0	1.0	0.0	0.0	0.0	

```
1 print("Head:\n",logistic_df.head(), "\n\n")
2
3 print("DTypes:\n",logistic_df.dtypes, "\n\n")
.
```

```

4
5 print("Description:\n",logistic_df.describe())

Head:
   Age  Number of sexual partners  First sexual intercourse \
0    18                        4.0                      15.0
1    15                        1.0                      14.0
2    34                        1.0                      NaN
3    52                        5.0                      16.0
4    46                        3.0                      21.0

   Num of pregnancies  Smokes  Smokes (years)  Smokes (packs/year) \
0                    1.0    0.0              0.0                0.0
1                    1.0    0.0              0.0                0.0
2                    1.0    0.0              0.0                0.0
3                    4.0    1.0             37.0               37.0
4                    4.0    0.0              0.0                0.0

   Hormonal Contraceptives  Hormonal Contraceptives (years)  IUD  ... \
0                      0.0                                0.0  0.0  ...
1                      0.0                                0.0  0.0  ...
2                      0.0                                0.0  0.0  ...
3                      1.0                                3.0  0.0  ...
4                      1.0                               15.0  0.0  ...

   STDs: Time since first diagnosis  STDs: Time since last diagnosis \
0                                NaN                                NaN
1                                NaN                                NaN
2                                NaN                                NaN
3                                NaN                                NaN
4                                NaN                                NaN

   Dx:Cancer  Dx:CIN  Dx:HPV  Dx  Hinselmann  Schiller  Citology  Biopsy
0           0       0       0  0           0         0         0       0
1           0       0       0  0           0         0         0       0
2           0       0       0  0           0         0         0       0
3           1       0       1  0           0         0         0       0
4           0       0       0  0           0         0         0       0

```

[5 rows x 36 columns]

```

DTypes:
Age                                int64
Number of sexual partners          float64
First sexual intercourse            float64
Num of pregnancies                  float64
Smokes                             float64
Smokes (years)                     float64
Smokes (packs/year)                float64
Hormonal Contraceptives            float64
Hormonal Contraceptives (years)    float64
IUD                                float64
IUD (years)                         float64
STDs                                float64
STDs (number)                       float64
STDs:condylomatosis                float64
STDs:cervical condylomatosis        float64
STDs:vaginal condylomatosis         float64
STDs:vulvo-perineal condylomatosis  float64
STDs:syphilis                       float64

```

Missing Values

```

1 print("Nulls:\n",logistic_df.isnull().sum())

Nulls:
Age                                0
Number of sexual partners          26
First sexual intercourse            7
Num of pregnancies                  56
Smokes                             13
Smokes (years)                     13
Smokes (packs/year)                13
Hormonal Contraceptives            108
Hormonal Contraceptives (years)    108
IUD                                117
IUD (years)                         117
STDs                                105
STDs (number)                       105
STDs:condylomatosis                105
STDs:cervical condylomatosis        105

```

```

STDs:vaginal condylomatosis      105
STDs:vulvo-perineal condylomatosis 105
STDs:syphilis                    105
STDs:pelvic inflammatory disease 105
STDs:genital herpes              105
STDs:molluscum contagiosum       105
STDs:AIDS                       105
STDs:HIV                        105
STDs:Hepatitis B                 105
STDs:HPV                        105
STDs: Number of diagnosis        0
STDs: Time since first diagnosis  787
STDs: Time since last diagnosis  787
Dx:Cancer                       0
Dx:CIN                          0
Dx:HPV                          0
Dx                              0
Hinselmann                     0
Schiller                        0
Citology                       0
Biopsy                         0
dtype: int64

```

```

1 print("Unique values in 'IUD' column:")
2 print(logistic_df['IUD'].unique())
3
4 print("\n\nUnique values in 'STDs:HIV ' column:")
5 print(logistic_df['STDs:HIV'].unique())

```

```

Unique values in 'IUD' column:
[ 0.  1. nan]

```

```

Unique values in 'STDs:HIV ' column:
[ 0.  1. nan]

```

Duplicates

```

1 duplicates = logistic_df.duplicated()
2 print("Duplicates:\n\n", duplicates, "\n\n")
3 print("Duplicate Rows:\n\n",logistic_df[duplicates])

```

Duplicates:

```

0      False
1      False
2      False
3      False
4      False
...
853     False
854     False
855     False
856     False
857     False
Length: 858, dtype: bool

```

Duplicate Rows:

	Age	Number of sexual partners	First sexual intercourse \
66	34	3.0	19.0
234	25	NaN	18.0
255	25	2.0	18.0
356	18	1.0	17.0
395	18	1.0	18.0
406	17	1.0	17.0
419	19	4.0	14.0
431	18	1.0	14.0
435	17	2.0	15.0
440	15	1.0	14.0
442	16	1.0	15.0
453	15	1.0	15.0
454	15	1.0	14.0
466	15	1.0	14.0
486	28	2.0	20.0
525	17	1.0	16.0
530	21	4.0	15.0
536	16	1.0	14.0
575	17	2.0	15.0

580	17	2.0	15.0
638	21	1.0	20.0
715	15	2.0	14.0
785	28	1.0	19.0

	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	\
66	3.0	0.0	0.0	0.0	
234	2.0	0.0	0.0	0.0	
255	2.0	0.0	0.0	0.0	
356	1.0	0.0	0.0	0.0	
395	1.0	0.0	0.0	0.0	
406	1.0	0.0	0.0	0.0	
419	1.0	0.0	0.0	0.0	
431	2.0	0.0	0.0	0.0	
435	1.0	0.0	0.0	0.0	
440	1.0	0.0	0.0	0.0	
442	1.0	0.0	0.0	0.0	
453	1.0	0.0	0.0	0.0	
454	1.0	0.0	0.0	0.0	
466	1.0	0.0	0.0	0.0	

Steps:

Cleaning Missing Values

Since 'STDs: Time since first diagnosis' and 'STDs: Time since last diagnosis' have a huge amount of missing values, accounting to about 92%, we will just drop these 2 columns to remain the integrity of the data

```
1 logistic_df.columns
```

```
Index(['Age', 'Number of sexual partners', 'First sexual intercourse',
      'Num of pregnancies', 'Smokes', 'Smokes (years)', 'Smokes (packs/year)',
      'Hormonal Contraceptives', 'Hormonal Contraceptives (years)', 'IUD',
      'IUD (years)', 'STDs', 'STDs (number)', 'STDs:condylomatosis',
      'STDs:cervical condylomatosis', 'STDs:vaginal condylomatosis',
      'STDs:vulvo-perineal condylomatosis', 'STDs:syphilis',
      'STDs:pelvic inflammatory disease', 'STDs:genital herpes',
      'STDs:molluscum contagiosum', 'STDs:AIDS', 'STDs:HIV',
      'STDs:Hepatitis B', 'STDs:HPV', 'STDs: Number of diagnosis',
      'STDs: Time since first diagnosis', 'STDs: Time since last diagnosis',
      'Dx:Cancer', 'Dx:CIN', 'Dx:HPV', 'Dx', 'Hinselmann', 'Schiller',
      'Citology', 'Biopsy'],
      dtype='object')
```

```
1 drop_columns = ['STDs: Time since first diagnosis', 'STDs: Time since last diagnosis']
2 logistic_df.drop(drop_columns, axis=1, inplace=True)
```

```
1 logistic_df.columns
```

```
Index(['Age', 'Number of sexual partners', 'First sexual intercourse',
      'Num of pregnancies', 'Smokes', 'Smokes (years)', 'Smokes (packs/year)',
      'Hormonal Contraceptives', 'Hormonal Contraceptives (years)', 'IUD',
      'IUD (years)', 'STDs', 'STDs (number)', 'STDs:condylomatosis',
      'STDs:cervical condylomatosis', 'STDs:vaginal condylomatosis',
      'STDs:vulvo-perineal condylomatosis', 'STDs:syphilis',
      'STDs:pelvic inflammatory disease', 'STDs:genital herpes',
      'STDs:molluscum contagiosum', 'STDs:AIDS', 'STDs:HIV',
      'STDs:Hepatitis B', 'STDs:HPV', 'STDs: Number of diagnosis',
      'Dx:Cancer', 'Dx:CIN', 'Dx:HPV', 'Dx', 'Hinselmann', 'Schiller',
      'Citology', 'Biopsy'],
      dtype='object')
```

```
1 no_missing_df = logistic_df.dropna()
```

```
1 no_missing_df
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs:HPV	STDs Number o diagnosi
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	0.0	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	0.0	
5	42	3.0	23.0	2.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
...	
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	0.0	
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	0.0	
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	0.0	
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	0.0	

✓ Cleaning Duplicates

```
1 new_df = no_missing_df.drop_duplicates()
```

```
1 new_df
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs:HPV	STDs Number o diagnosi
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	0.0	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	0.0	
5	42	3.0	23.0	2.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
...	
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	0.0	
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	0.0	
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	0.0	
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	0.0	
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	0.0	

✓ Checking:

```
1 print("Nulls:\n",new_df.isnull().sum())
```

```
Nulls:
Age                0
Number of sexual partners  0
First sexual intercourse  0
Num of pregnancies    0
Smokes              0
Smokes (years)      0
Smokes (packs/year)  0
Hormonal Contraceptives  0
Hormonal Contraceptives (years)  0
IUD                 0
IUD (years)         0
STDs                0
STDs (number)       0
```

```

STDs:condylomatosis      0
STDs:cervical condylomatosis  0
STDs:vaginal condylomatosis  0
STDs:vulvo-perineal condylomatosis  0
STDs:syphilis            0
STDs:pelvic inflammatory disease  0
STDs:genital herpes      0
STDs:molluscum contagiosum  0
STDs:AIDS                 0
STDs:HIV                  0
STDs:Hepatitis B         0
STDs:HPV                  0
STDs: Number of diagnosis  0
Dx:Cancer                 0
Dx:CIN                    0
Dx:HPV                    0
Dx                        0
Hinselmann                0
Schiller                  0
Citology                  0
Biopsy                    0
dtype: int64

```

```

1 print("Duplicates:\n\n", duplicates, "\n\n")
2 print("Duplicate Rows:\n\n",new_df[duplicates])

```

Duplicates:

```

0      False
1      False
2      False
3      False
4      False
...
853    False
854    False
855    False
856    False
857    False
Length: 858, dtype: bool

```

Duplicate Rows:

```

Empty DataFrame
Columns: [Age, Number of sexual partners, First sexual intercourse, Num of pregnancies, Smokes, Smokes (years), Smokes (packs/year), Hor
Index: []

```

```

[0 rows x 34 columns]
<ipython-input-119-0d93af64aaf1>:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
  print("Duplicate Rows:\n\n",new_df[duplicates])

```

```

1 original_data = len(logistic_df)
2 new_data = len(new_df)
3 rows_dropped = original_data - new_data
4 print("Number of rows dropped: ", rows_dropped)

```

Number of rows dropped: 208

Great! We have removed all the missing and duplicated values! Let's proceed to EDA

✓ Task: Find the relationship between Age and Cervical Cancer

```

1 print("Unique values in 'Age' column:")
2 print(new_df['Age'].unique())

```

```

Unique values in 'Age' column:
[18 15 52 46 42 51 26 45 44 27 43 40 41 39 37 38 36 35 33 34 31 32 30 23
 28 29 25 21 24 22 20 48 19 17 16 14 79 84 47 13 70 50 49]

```

```


1 filtered_new_df = new_df[['Age', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Age')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))

```

```

6 plt.bar(grouped_new_df_sorted['Age'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Age')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by Age')
10 plt.show()

```



We can see that the number of people with cervic cancer are only present in 20s - 50s

✓ **Task: Find the relationship between people that has a high number of sexual partners to cervical cancer**

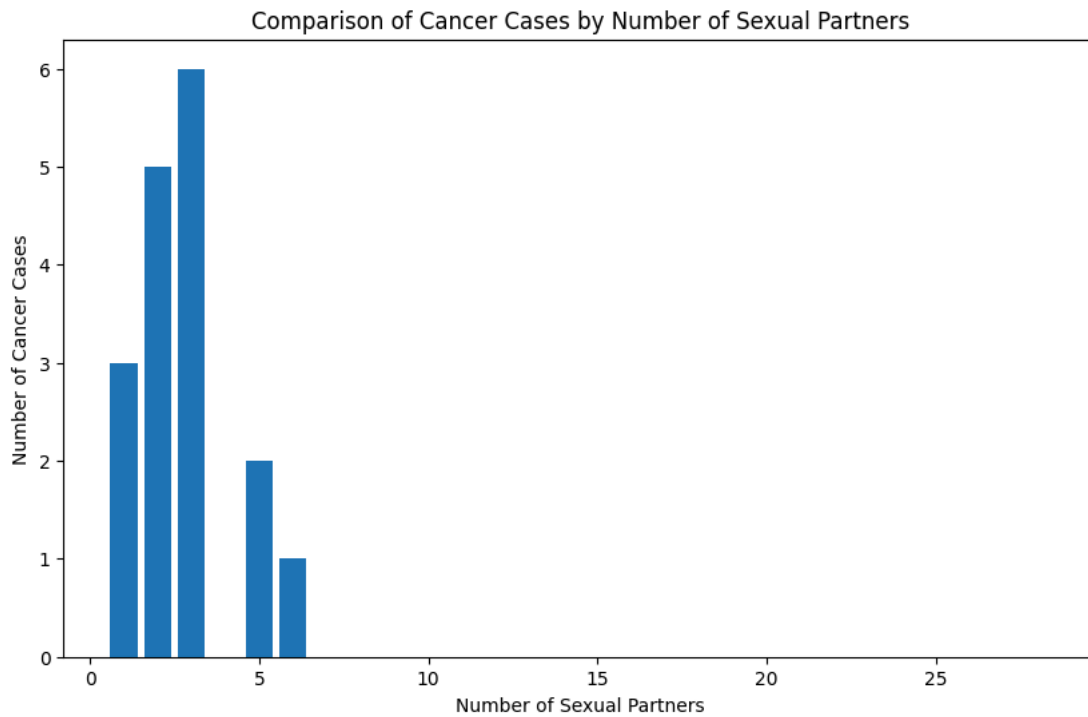
```

1 print("Unique values in 'Number of sexual partners' column:")
2 print(new_df['Number of sexual partners'].unique())

Unique values in 'Number of sexual partners' column:
[ 4.  1.  5.  3.  2.  6.  8.  7. 28.]

1 filtered_new_df = new_df[['Number of sexual partners', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Number of sexual partners')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Number of sexual partners'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Number of Sexual Partners')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by Number of Sexual Partners')
10 plt.show()

```

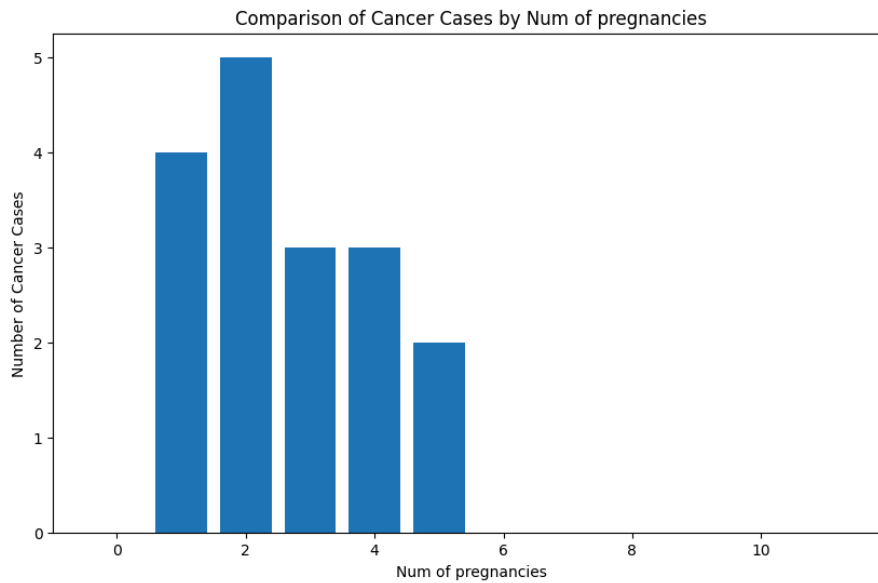
Conclusion, it seems that a high number of sexual partners do not correlate with having a cervic cancer, as there are not even a single cancer case for those with more than 7 up to 28 sexual partners

✓ Task: Find the relationship between pregnancies to cervical cancer

```
1 print("Unique values in 'Num of pregnancies' column:")
2 print(new_df['Num of pregnancies'].unique())
```

```
Unique values in 'Num of pregnancies' column:
[ 1.  4.  2.  6.  3.  5.  8.  7.  0. 11. 10.]
```

```
1 filtered_new_df = new_df[['Num of pregnancies', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Num of pregnancies')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Num of pregnancies'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Num of pregnancies')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by Num of pregnancies')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Num of pregnancies'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and Num of pregnancies:", correlation_coefficient)
```

Correlation between cervix cancer and Num of pregnancies: 0.03283919105903373


Same with the number of sexual partners, a higher amount of pregnancies does not seem to correlate to higher cases of people with cervic cancer

✓ Task: Find the relationship between smoking to cervical cancer

```
1 print("Unique values in 'Smokes' column:")
2 print(new_df['Smokes'].unique())
```

Unique values in 'Smokes' column:
[0. 1.]

```
1 filtered_new_df = new_df[['Smokes', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Smokes')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Smokes'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Smokes')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by Smoking')
10 plt.show()
```



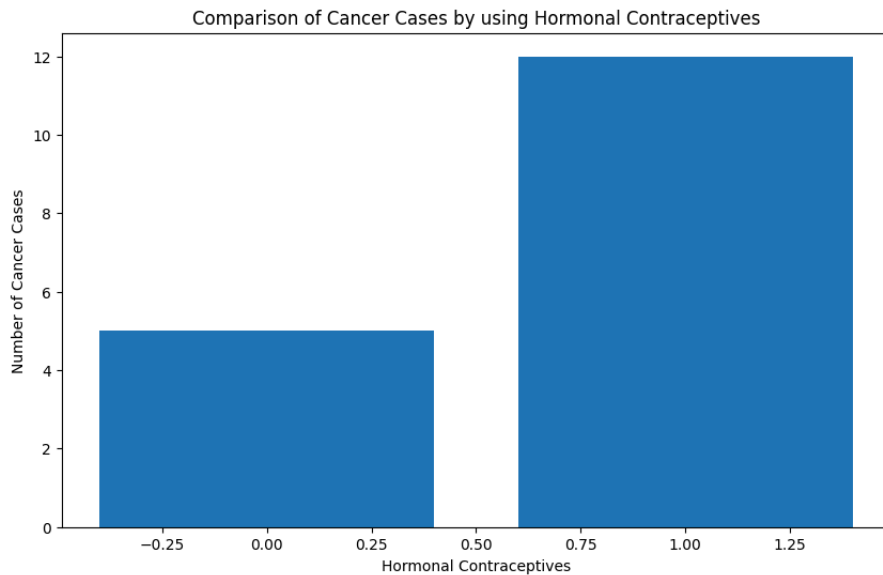
From the plot above, we can see that smoking does not affect whether the person has cervic cancer, as there are more cancer cases in those that do not smoke

✓ Task: Find the relationship between using Hormonal Contraceptives to cervical cancer

```
1 print("Unique values in 'Hormonal Contraceptives' column:")
2 print(new_df['Hormonal Contraceptives'].unique())
```

```
Unique values in 'Hormonal Contraceptives' column:
[0. 1.]
```

```
1 filtered_new_df = new_df[['Hormonal Contraceptives', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Hormonal Contraceptives')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Hormonal Contraceptives'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Hormonal Contraceptives')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by using Hormonal Contraceptives')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Hormonal Contraceptives'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and usage of Hormonal Contraceptives:", correlation_coefficient)

Correlation between cervix cancer and usage of Hormonal Contraceptives: 0.017413219501020982
```

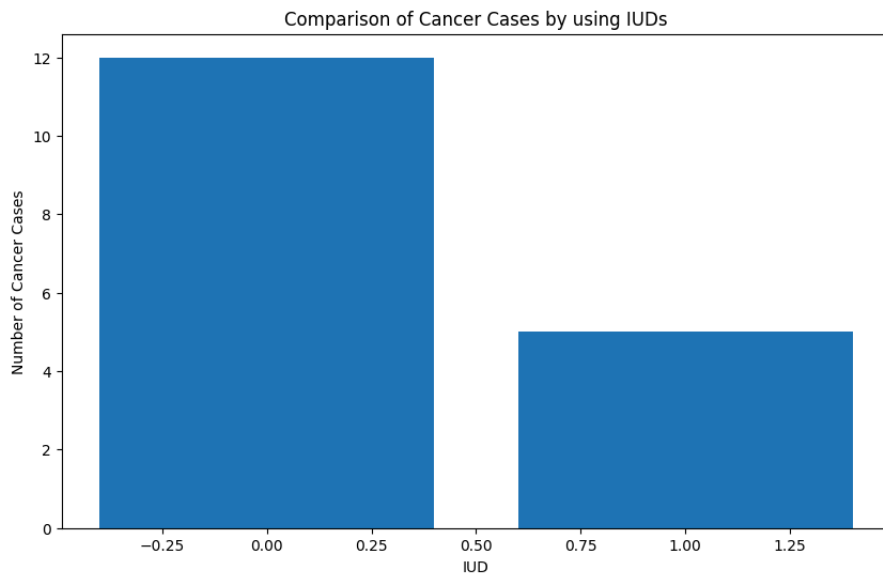
We can see that those that use hormonal contraceptives have higher cancer cases than those that do not.

✓ Task: Find the relationship between using IUD to cervical cancer

```
1 print("Unique values in 'IUD' column:")
2 print(new_df['IUD'].unique())

Unique values in 'IUD' column:
[0. 1.]

1 filtered_new_df = new_df[['IUD', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('IUD')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['IUD'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('IUD')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by using IUDs')
10 plt.show()
```



There is still a significant difference as more than half of those that do not use IUDs have cancer compared to those that uses them, therefore, we can infer that using IUDs do not affect whether an individual has cervix cancer or not

```
1 correlation_coefficient = new_df['IUD'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and usage of IUD:", correlation_coefficient)
```

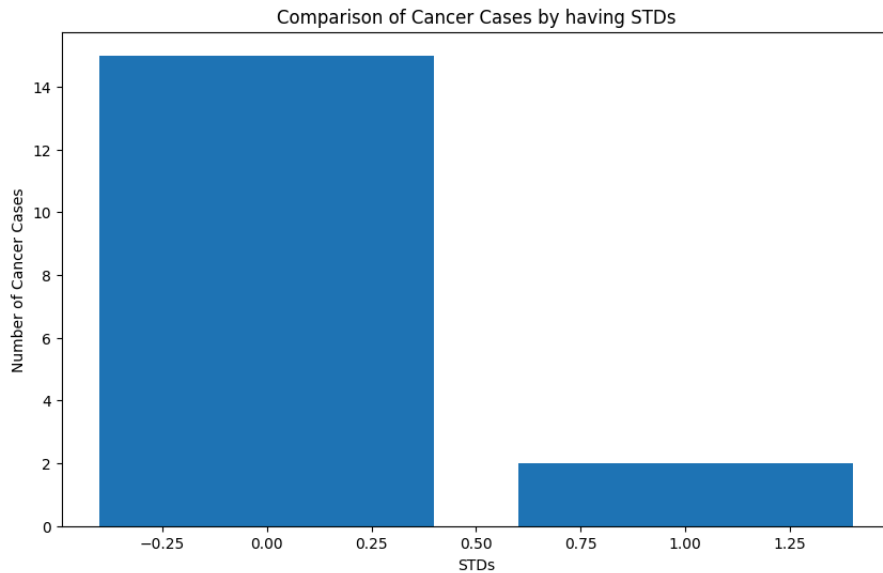
Correlation between cervix cancer and usage of IUD: 0.09168034316657442

✓ Task: Find the relationship between having STDs to cervical cancer

```
1 print("Unique values in 'STDs' column:")
2 print(new_df['STDs'].unique())
```

Unique values in 'STDs' column:
[0. 1.]

```
1 filtered_new_df = new_df[['STDs', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('STDs')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['STDs'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('STDs')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by having STDs')
10 plt.show()
```



```
1 correlation_coefficient = new_df['STDs'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and STDs:", correlation_coefficient)

Correlation between cervix cancer and STDs: 0.009639925649468188
```


Surprisingly, the number of cancer cases is higher for those that do not have STD than those that has them. Along with the correlation analysis, we can say that having an STD doesn't mean that they will also have cervix cancer

✓ Task: Find the relationship between diagnosed of CIN to cervical cancer

```
1 print("Unique values in 'Dx:CIN' column:")
2 print(new_df['Dx:CIN'].unique())

Unique values in 'Dx:CIN' column:
[0 1]

1 filtered_new_df = new_df[['Dx:CIN', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Dx:CIN')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Dx:CIN'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Dx:CIN')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by being diagnosed of Cervical intraepithelial neoplasia(CIN)')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Dx:CIN'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and Dx:CIN:", correlation_coefficient)
```

Correlation between cervix cancer and Dx:CIN: -0.011159149417981603


A negative correlation combined with the graph tells us that the people with this diagnosis absolutely has nothing to do with having a cervix cancer.

✓ Task: Find the relationship between diagnosed of HPV to cervical cancer

```
1 print("Unique values in 'Dx:HPV' column:")
2 print(new_df['Dx:HPV'].unique())
```

Unique values in 'Dx:HPV' column:
[0 1]

```
1 filtered_new_df = new_df[['Dx:HPV', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Dx:HPV')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Dx:HPV'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Dx:HPV')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases by being diagnosed of Human papillomavirus (HPV)')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Dx:HPV'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and Dx:HPV:", correlation_coefficient)
```

Correlation between cervix cancer and Dx:HPV: 0.9071639187403662

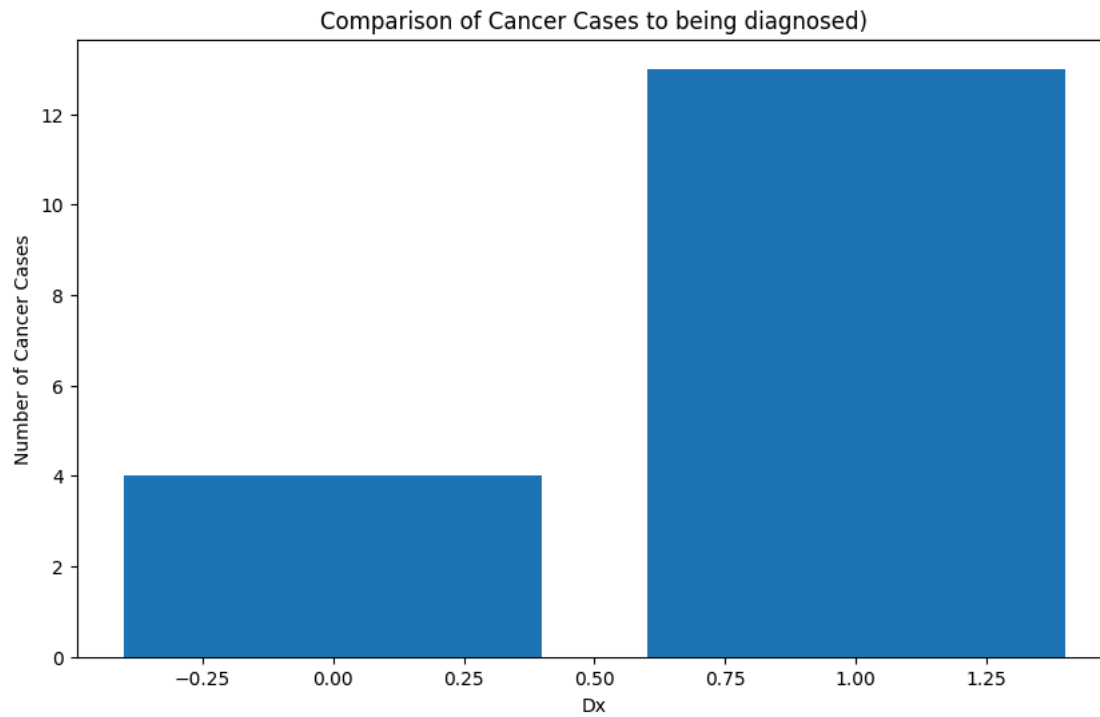
The comparison plus the correlation of HPV to Cervix Cancer tells us that there is a strong relationship between the two.

✓ Task: Find the relationship between doing a dx test to cervical cancer

```
1 print("Unique values in 'Dx column:")
2 print(new_df['Dx'].unique())
```

Unique values in 'Dx column:
[0 1]

```
1 filtered_new_df = new_df[['Dx', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Dx')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Dx'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Dx')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases to doing the OncoType dx test')
10 plt.show()
```

```
1 correlation_coefficient = new_df['Dx'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and doing the onco type dx test:", correlation_coefficient)

Correlation between cervix cancer and doing the onco type dx test: 0.7827375530131594
```

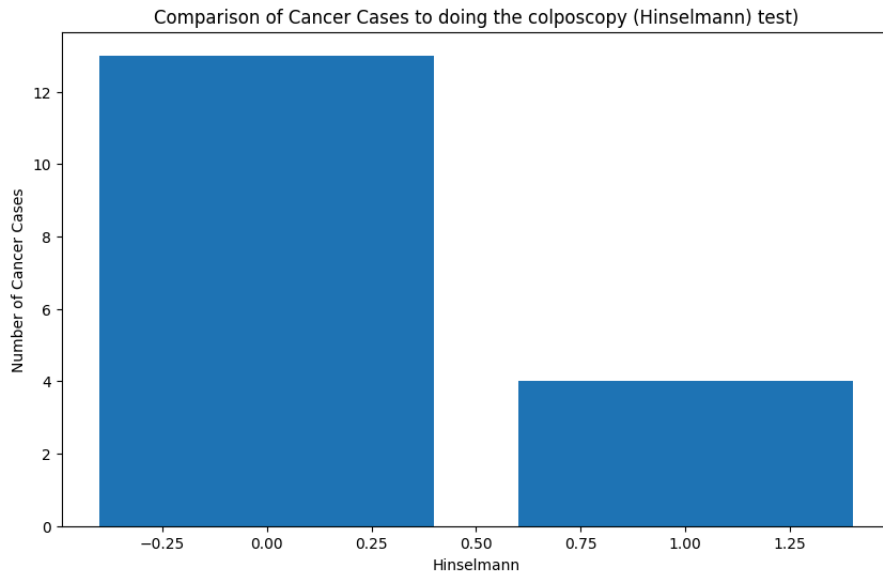
From the comparison and correlation above, we can infer that those that do the test are found to have cancer more than those that do not

✓ Task: Find the relationship between doing colposcopy to cervical cancer

```
1 print("Unique values in 'Hinselmann' column:")
2 print(new_df['Hinselmann'].unique())

Unique values in 'Hinselmann' column:
[0 1]

1 filtered_new_df = new_df[['Hinselmann', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Hinselmann')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Hinselmann'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Hinselmann')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases to doing the colposcopy (Hinselmann) test')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Hinselmann'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and doing the colposcopy test:", correlation_coefficient)
```

Correlation between cervix cancer and doing the colposcopy test: 0.1477282313363026


There is only a weak correlation between the colposcopy test and having cancer, on top of the huge difference between those that do not do the test having cancer than those that did it.

✓ Task: Find the relationship between doing the Schiller test to having cervical cancer

```
1 print("Unique values in 'Schiller' column:")
2 print(new_df['Schiller'].unique())
```

Unique values in 'Schiller' column:
[0 1]

```
1 filtered_new_df = new_df[['Schiller', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Schiller')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Schiller'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Schiller')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases to doing the Schiller test')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Schiller'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and doing the Schiller test:", correlation_coefficient)

Correlation between cervix cancer and doing the Schiller test: 0.14368918194911728
```

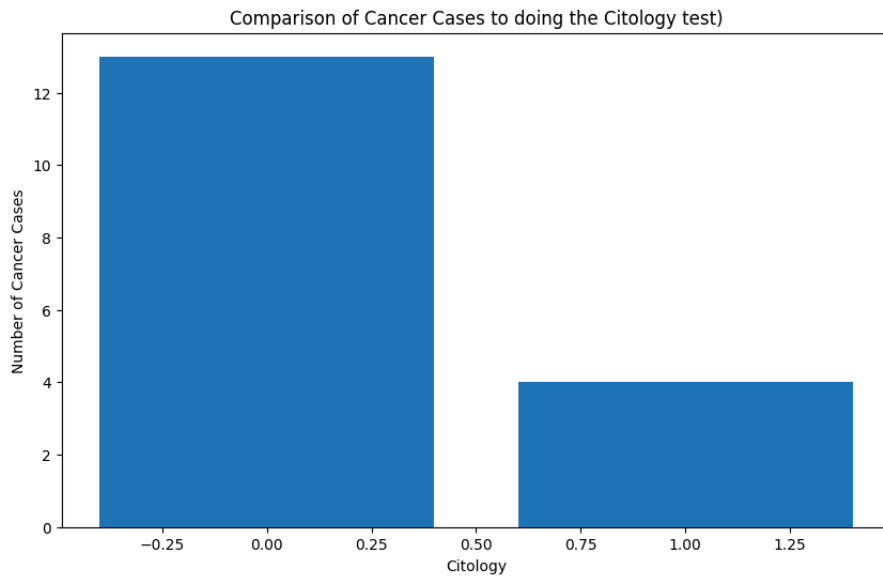
While the plot shows that there isn't that huge of a difference between those that took the schiller test to having cancer or not, the correlation coefficient says otherwise.

✓ Task: Find the relationship between doing the Citology test to having cervical cancer

```
1 print("Unique values in 'Citology column:")
2 print(new_df['Citology'].unique())

Unique values in 'Citology column:
[0 1]

1 filtered_new_df = new_df[['Citology', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Citology')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Citology'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Citology')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases to doing the Citology test')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Citology'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and doing the Citology test:", correlation_coefficient)

Correlation between cervix cancer and doing the Citology test: 0.1235181374669367
```

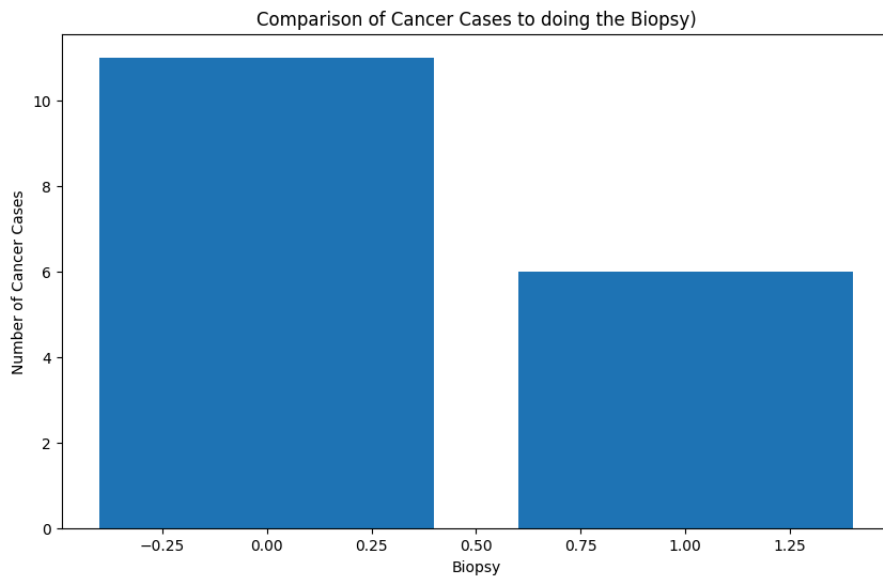
Same with the previous one, there isn't a relationship between the Citology test to having a cervical cancer

✓ Task: Find the relationship between doing the Biopsy to having cervical cancer

```
1 print("Unique values in 'Biopsy column:")
2 print(new_df['Biopsy'].unique())

Unique values in 'Biopsy column:
[0 1]

1 filtered_new_df = new_df[['Biopsy', 'Dx:Cancer']].copy()
2 grouped_new_df = filtered_new_df.groupby('Biopsy')['Dx:Cancer'].sum().reset_index()
3 grouped_new_df_sorted = grouped_new_df.sort_values(by='Dx:Cancer', ascending=False)
4
5 plt.figure(figsize=(10, 6))
6 plt.bar(grouped_new_df_sorted['Biopsy'], grouped_new_df_sorted['Dx:Cancer'])
7 plt.xlabel('Biopsy')
8 plt.ylabel('Number of Cancer Cases')
9 plt.title('Comparison of Cancer Cases to doing the Biopsy')
10 plt.show()
```



```
1 correlation_coefficient = new_df['Biopsy'].corr(new_df['Dx:Cancer'])
2 print("Correlation between cervix cancer and doing the Biopsy:", correlation_coefficient)
```

Correlation between cervix cancer and doing the Biopsy: 0.1860789482965945

From the plot and coefficient above, we can see that there isn't a relationship between doing a biopsy to having a cervical cancer.

We can then remove the columns relating to sexual partners, pregnancies, smoking, IUDs, STDs, Dx:CIN, Hinselmann, Schiller, Citology, and biopsy. as they do not affect whether an individual will have cervix cancer or not

Thus, we will only use Age, Dx:HPV, and Dx as the Features and Dx:Cancer as the target.

✓ Logistic Regression

✓ Declare Feature Vector and Target Variable

```
1 X = new_df[['Age', 'Dx:HPV', 'Dx']]
2 y = new_df['Dx:Cancer']

1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0, stratify=y)
2 X_train.shape, X_test.shape

((520, 3), (130, 3))

1 print("Shape of X_train:", X_train.shape)
2 print("Shape of X_test:", X_test.shape)
3 print("Shape of y_train:", y_train.shape)
4 print("Shape of y_test:", y_test.shape)

Shape of X_train: (520, 3)
Shape of X_test: (130, 3)
Shape of y_train: (520,)
Shape of y_test: (130,)
```

✓ Feature Engineering

```
1 categorical = [col for col in X_train.columns if X_train[col].dtypes == '0']
2 numerical = [col for col in X_train.columns if X_train[col].dtypes != '0']
3 print("X_train Dtypes:\n", X_train.dtypes)
4 print("\nCategorical columns:\n", categorical)
5 print("\nNumerical columns:\n", numerical)
```

```
X_train Dtypes:
Age      int64
Dx:HPV   int64
Dx       int64
dtype: object

Categorical columns:
[]

Numerical columns:
['Age', 'Dx:HPV', 'Dx']
```

✓ Engineering Missing Values in Numerical Variables

```
1 print("Missing Values in the numerical variables of X_train:\n", X_train[numerical].isnull().sum())
2 print("\nMissing Values in the numerical variables of X_test:\n", X_test[numerical].isnull().sum())
```

```
Missing Values in the numerical variables of X_train:
Age      0
Dx:HPV   0
Dx       0
dtype: int64

Missing Values in the numerical variables of X_test:
Age      0
Dx:HPV   0
Dx       0
dtype: int64
```

```
1 #Print Percentage of missing values in the numerical variables in training set
2 for col in numerical:
3     if X_train[col].isnull().mean()>0:
4         print(col, round(X_train[col].isnull().mean(), 4))
```

✓ Feature Scaling

```
1 print("X_train description:\n", X_train.describe())
2 cols = X_train.columns
3 scaler = MinMaxScaler()
4 X_train = scaler.fit_transform(X_train)
5 X_test = scaler.transform(X_test)
```

```
X_train description:
      Age      Dx:HPV      Dx
count  520.000000  520.000000  520.000000
mean    27.386538    0.025000    0.026923
std      8.711752    0.156275    0.162015
min     13.000000    0.000000    0.000000
25%     21.000000    0.000000    0.000000
50%     26.000000    0.000000    0.000000
75%     33.000000    0.000000    0.000000
max     84.000000    1.000000    1.000000
```

```
1 X_train = pd.DataFrame(X_train, columns = [cols])
2 X_test = pd.DataFrame(X_test, columns = [cols])
3 print("X_train description:\n", X_train.describe())
```

```
X_train description:
      Age      Dx:HPV      Dx
count  520.000000  520.000000  520.000000
mean     0.202627    0.025000    0.026923
std      0.122701    0.156275    0.162015
min      0.000000    0.000000    0.000000
```

25%	0.112676	0.000000	0.000000
50%	0.183099	0.000000	0.000000
75%	0.281690	0.000000	0.000000
max	1.000000	1.000000	1.000000

- Model Training

```
1 logreg = LogisticRegression(solver = 'liblinear', random_state = 0)
2 logreg.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression(random_state=0, solver='liblinear')
```

- ▼ Predict Results

```
1 y_pred_test = logreg.predict(X_test)
2 y_pred_test
```

```
array([[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
1 #Probability of getting output as 0
2 logreg.predict_proba(X_test)[:,:0]
```

```
array([0.98370138, 0.98324086, 0.98294117, 0.98353538, 0.98328325,
       0.98378377, 0.98302733, 0.98366004, 0.98302733, 0.9829843 ,
       0.98349362, 0.98361859, 0.98285458, 0.9838248 , 0.98345175,
       0.98324086, 0.98311306, 0.98315577, 0.98398794, 0.98294117,
       0.98390657, 0.98336771, 0.70527716, 0.98302733, 0.15812346,
       0.9829843 , 0.9838248 , 0.98315577, 0.98340978, 0.9838248 ,
       0.98311306, 0.98294117, 0.98307025, 0.98332553, 0.98289793,
       0.98319837, 0.98340978, 0.98430936, 0.98311306, 0.98319837,
       0.98289793, 0.98332553, 0.98328325, 0.98366004, 0.98366004,
       0.98357703, 0.98315577, 0.98336771, 0.98332553, 0.98311306,
       0.98328325, 0.9829843 , 0.98345175, 0.98311306, 0.98345175,
       0.98361859, 0.98319837, 0.98353538, 0.98361859, 0.98357703,
       0.98307025, 0.98319837, 0.98332553, 0.9829843 , 0.98285458,
       0.98307025, 0.98289793, 0.98319837, 0.98345175, 0.98311306,
       0.98328325, 0.98386574, 0.98340978, 0.98267555, 0.98315577,
       0.98324086, 0.98311306, 0.98340978, 0.98357703, 0.98349362,
       0.98340978, 0.98294117, 0.98311306, 0.98370138, 0.98285458,
       0.98285458, 0.98294117, 0.98345175, 0.98353538, 0.98294117,
       0.98340978, 0.98386574, 0.98311306, 0.98311306, 0.98302733,
       0.98315577, 0.98285458, 0.98336771, 0.9829843 , 0.98328325,
       0.98324086, 0.9829843 , 0.98285458, 0.98357703, 0.98315577,
       0.15915478, 0.98324086, 0.98357703, 0.98302733, 0.98349362,
       0.98361859, 0.98353538, 0.98366004, 0.98402847, 0.98361859,
       0.98353538, 0.98328325, 0.9838248 , 0.98289793, 0.98340978,
       0.98319837, 0.98311306, 0.98311306, 0.98289793, 0.98534799,
       0.98289793, 0.98366004, 0.98324086, 0.98281112, 0.98302733])
```

```
1 #Probability of getting output as 1
2 logreg.predict_proba(X_test)[: ,1]
```

```
array([0.01629862, 0.01675914, 0.01705883, 0.01646462, 0.01671675,
       0.01621623, 0.01697267, 0.01633996, 0.01697267, 0.0170157,
       0.01659638, 0.01638141, 0.01714542, 0.0161752, 0.01654825,
       0.01675914, 0.01688694, 0.01684423, 0.01601206, 0.01705883,
       0.01609343, 0.01663229, 0.29472284, 0.01697267, 0.84187654,
       0.0170157, 0.0161752, 0.01684423, 0.01659022, 0.0161752,
       0.01688694, 0.01705883, 0.01692975, 0.01667447, 0.01710207,
       0.01680163, 0.01659022, 0.01659064, 0.01688694, 0.01680163,
       0.01710207, 0.01667447, 0.01671675, 0.01633996, 0.01633996,
       0.01642297, 0.01684423, 0.01663229, 0.01667447, 0.01688694,
       0.01671675, 0.0170157, 0.01654825, 0.01688694, 0.01654825,
       0.01638141, 0.01680163, 0.01646462, 0.01638141, 0.01642297,
       0.01692975, 0.01680163, 0.01667447, 0.0170157, 0.01714542,
       0.01692975, 0.01710207, 0.01680163, 0.01654825, 0.01688694,
       0.01671675, 0.01613426, 0.01659022, 0.01723245, 0.01684423,
       0.01675914, 0.01688694, 0.01659022, 0.01642297, 0.01659638,
```

- ✓ Check Accuracy Score

A high model accuracy score!

- ✓ Compare the train-set and test-set accuracy

```
1 new_df['Dx:Cancer'].value_counts()
```

```
Dx:Cancer
0      633
1       17
Name: count, dtype: int64
```

```
1 y_pred_train = logreg.predict(X_train)
2 print("Prediction Training:\n", y_pred_train)
3 print('\nTraining-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

[illegible]

Training-set accuracy score: 0.9904

- ✓ Check for overfitting and underfitting

```
1 print('Training set score: {:.4f}'.format(logreg.score(X_train, y_train)))
2 print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))
```

```
Training set score: 0.9904
Test set score: 0.9923
```

As the training set score and test set score are quite comparable to each other, there is no question of overfitting

- ✓ Using $C = 100$

```
1 logreg100 = LogisticRegression(C = 100, solver = 'liblinear', random_state = 0)
2 logreg100.fit(X_train, y_train)
```



```

▼ LogisticRegression
LogisticRegression(C=100, random_state=0, solver='liblinear')

```

```

1 print('Training set score: {:.4f}'.format(logreg100.score(X_train, y_train)))
2 print('Test set score: {:.4f}'.format(logreg100.score(X_test, y_test)))

```

```

Training set score: 0.9962
Test set score: 1.0000

```

We can see that C = 100 results in higher accuracy, even making the test score flat 1 as well, meaning that this model performed the best and no question of overfitting or underfitting

▼ Using C = 0.01

```

1 logreg001 = LogisticRegression(C = 0.01, solver = 'liblinear', random_state = 0)
2 logreg001.fit(X_train, y_train)

```

```

▼ LogisticRegression
LogisticRegression(C=0.01, random_state=0, solver='liblinear')

```

```

1 print('Training set score: {:.4f}'.format(logreg001.score(X_train, y_train)))
2 print('Test set score: {:.4f}'.format(logreg001.score(X_test, y_test)))

```

```

Training set score: 0.9731
Test set score: 0.9769

```

As we can see, the scores remain to be almost the same as each other while remaining to be high even if c = 0.01. There is now no question for underfitting

▼ Comparing the model accuracy with null accuracy

```

1 y_test.value_counts()

```

```

Dx:Cancer
0      127
1         3
Name: count, dtype: int64

```

```

1 null_accuracy = (127/(127+3))
2 print('\nNull Accuracy Score: {0:0.4f}'.format(null_accuracy))

```

```

Null Accuracy Score: 0.9769

```

▼ Confusion Matrix

```

1 cm = confusion_matrix(y_test, y_pred_test)
2 print('Confusion matrix\n\n', cm)
3 print("\nTrue Positives(TP) = ", cm[0,0])
4 print("\nTrue Negatives(TN) = ", cm[1,1])
5 print("\nFalse Positives(FP) = ", cm[0,1])
6 print("\nFalse Negatives(FN) = ", cm[1,0])

```

```

Confusion matrix

```

```

[[127  0]
 [ 1  2]]

```

```

True Positives(TP) = 127

```

```

True Negatives(TN) = 2

```

```

False Positives(FP) = 0

```

```

False Negatives(FN) = 1

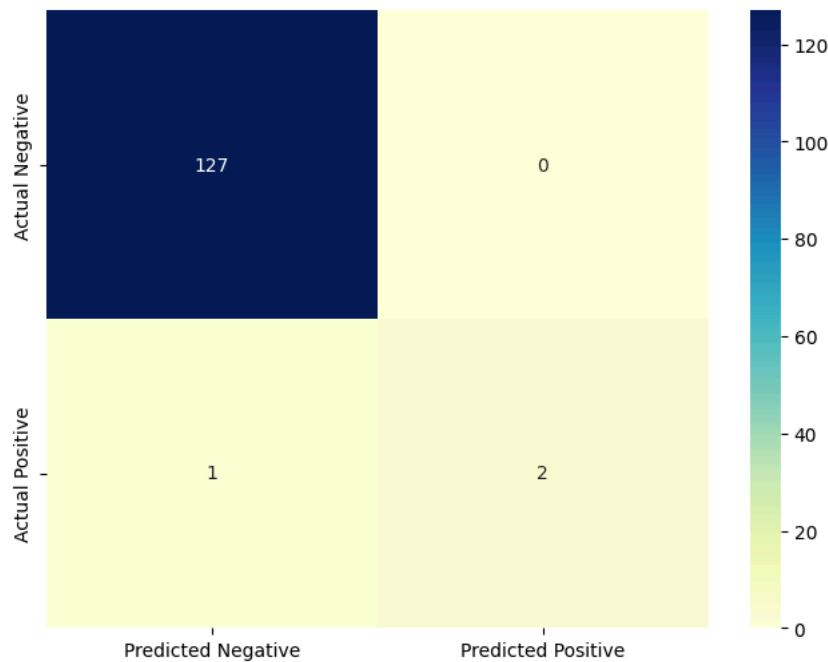
```

```

1 cm_matrix = pd.DataFrame(data=cm, columns=['Predicted Negative', 'Predicted Positive'],
2                           index=['Actual Negative', 'Actual Positive'])
3 plt.figure(figsize=(8, 6))
4 sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')

```

<Axes: >



```

1 from sklearn.metrics import roc_auc_score
2 precision = (127) / (127 + 0)
3 recall = (127) / (127 + 1)
4 f1 = 2 * ((precision * recall) / (precision + recall))
5 roc_auc = roc_auc_score(y_test, y_pred_test)
6
7 print(f'Precision: {precision:.4f}')
8 print(f'Recall: {recall:.4f}')
9 print(f'F1-score: {f1:.4f}')
10 print(f'ROC AUC: {roc_auc:.4f}')

```

```

Precision: 1.0000
Recall: 0.9922
F1-score: 0.9961
ROC AUC: 0.8333

```

Conclusion: The model I created had a high accuracy, especially in getting positive cases right even though they are severely overwhelmed by the number of negative cases, hardly ever making false positives. My scores in precision, recall, f1, and roc auc are also pretty high, meaning that my model is performing very well in terms of precision and its ability to distinguish the two cases. Therefore, I can say that I have developed a model that can accurately predict cervical cancer with the presence of the features at hand.