

Name: Garcia, John Carlos M.  
Section: CPE22S3  
Date Performed: 04/22/2024  
Date Submitted: 04/27/2024

## Logistic Regression Analysis

Task: Determine the origin of wines

### Setup

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split
6
7 %matplotlib inline

1 pip install ucimlrepo

Collecting ucimlrepo
  Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.6

1 from ucimlrepo import fetch_ucirepo
2
3 # fetch dataset
4 wine = fetch_ucirepo(id=109)
5
6 # data (as pandas dataframes)
7 X = wine.data.features
8 y = wine.data.targets
9
10 # metadata
11 print("Metadata:\n",wine.metadata)
12
13 # variable information
14 print("\nVariables:\n",wine.variables)
15

Metadata:
{'uci_id': 109, 'name': 'Wine', 'repository_url': 'https://archive.ics.uci.edu/dataset/109/wine', 'data_url': 'https://archive.ics.uci.edu/dataset/109/wine'}

Variables:
  name      role      type demographic \
0      class  Target  Categorical      None
1    Alcohol  Feature  Continuous      None
2  Malicacid  Feature  Continuous      None
3      Ash    Feature  Continuous      None
4  Alcalinity_of_ash  Feature  Continuous      None
5    Magnesium  Feature    Integer      None
6  Total_phenols  Feature  Continuous      None
7  Flavanoids    Feature  Continuous      None
8  Nonflavanoid_phenols  Feature  Continuous      None
9  Proanthocyanins  Feature  Continuous      None
10  Color_intensity  Feature  Continuous      None
11      Hue    Feature  Continuous      None
12  0D280_0D315_of_diluted_wines  Feature  Continuous      None
13      Proline  Feature    Integer      None

description units missing_values
0      None      None      no
1      None      None      no
```

2	None	None	no
3	None	None	no
4	None	None	no
5	None	None	no
6	None	None	no
7	None	None	no
8	None	None	no
9	None	None	no
10	None	None	no
11	None	None	no
12	None	None	no
13	None	None	no

## Concatenation

```
1 logistic_df = pd.concat([X, y], axis=1) #Combine both dataframes into one for more efficient manipulation of data
```

## Exploration

```
1 print("Head:\n",logistic_df.head(), "\n\n")
2
3 print("DTypes:\n",logistic_df.dtypes, "\n\n")
4
5 print("Description:\n",logistic_df.describe())
```

Head:

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	\
0	14.23	1.71	2.43	15.6	127	2.80	
1	13.20	1.78	2.14	11.2	100	2.65	
2	13.16	2.36	2.67	18.6	101	2.80	
3	14.37	1.95	2.50	16.8	113	3.85	
4	13.24	2.59	2.87	21.0	118	2.80	

	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	Color_intensity	Hue	\
0	3.06		0.28	2.29	5.64	1.04
1	2.76		0.26	1.28	4.38	1.05
2	3.24		0.30	2.81	5.68	1.03
3	3.49		0.24	2.18	7.80	0.86
4	2.69		0.39	1.82	4.32	1.04

	0D280_0D315_of_diluted_wines	Proline	class
0	3.92	1065	1
1	3.40	1050	1
2	3.17	1185	1
3	3.45	1480	1
4	2.93	735	1

DTypes:

Alcohol	float64
Malicacid	float64
Ash	float64
Alcalinity_of_ash	float64
Magnesium	int64
Total_phenols	float64
Flavanoids	float64
Nonflavanoid_phenols	float64
Proanthocyanins	float64
Color_intensity	float64
Hue	float64
0D280_0D315_of_diluted_wines	float64
Proline	int64
class	int64
dtype:	object

Description:

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	\
count	178.000000	178.000000	178.000000	178.000000	178.000000	
mean	13.000618	2.336348	2.366517	19.494944	99.741573	
std	0.811827	1.117146	0.274344	3.339564	14.282484	
min	11.030000	0.740000	1.360000	10.600000	70.000000	
25%	12.362500	1.602500	2.210000	17.200000	88.000000	
50%	13.050000	1.865000	2.360000	19.500000	98.000000	
75%	13.677500	3.082500	2.557500	21.500000	107.000000	
max	14.830000	5.800000	3.230000	30.000000	162.000000	

	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	\
count	178.000000	178.000000	178.000000	178.000000	
mean	2.295112	2.029270	0.361854	1.590899	
std	0.625851	0.998859	0.124453	0.572359	
min	0.980000	0.340000	0.130000	0.410000	
75%	1.717500	1.205000	0.270000	1.250000	

## Identification of Missing Values

```
1 print("Nulls:\n",logistic_df.isnull().sum())
```

```
Nulls:
Alcohol      0
Malicacid    0
Ash          0
Alcalinity_of_ash  0
Magnesium    0
Total_phenols  0
Flavanoids   0
Nonflavanoid_phenols  0
Proanthocyanins  0
Color_intensity  0
Hue          0
0D280_0D315_of_diluted_wines  0
Proline      0
class        0
dtype: int64
```

## Identification of Duplicated Rows

```
1 duplicates = logistic_df.duplicated()
2 print("Duplicates:\n\n", duplicates, "\n\n")
3 print("Duplicate Rows:\n\n",logistic_df[duplicates])
```

Duplicates:

```
0      False
1      False
2      False
3      False
4      False
...
173     False
174     False
175     False
176     False
177     False
Length: 178, dtype: bool
```

Duplicate Rows:

```
Empty DataFrame
Columns: [Alcohol, Malicacid, Ash, Alcalinity_of_ash, Magnesium, Total_phenols, Flavanoids, Nonflavanoid_phenols, Proanthocyanins, Color
Index: []
```

```
1 print("Columns:\n",logistic_df.columns, "\n\n")
2 print("\n\nUnique values in 'class' column:")
3 print(logistic_df['class'].unique())
4
5 print("\n\nUnique values in 'Alcohol' column:")
6 print(logistic_df['Alcohol'].unique())
7
8 print("\n\nUnique values in 'Malicacid' column:")
9 print(logistic_df['Malicacid'].unique())
10
11 print("\n\nUnique values in 'Ash' column:")
12 print(logistic_df['Ash'].unique())
13
14 print("\n\nUnique values in '0D280_0D315_of_diluted_wines' column:")
15 print(logistic_df['0D280_0D315_of_diluted_wines'].unique())
16
17 print("\n\nUnique values in 'Proline' column:")
18 print(logistic_df['Proline'].unique())
```

Unique values in 'Alcohol' column:

```
[14.23 13.2 13.16 14.37 13.24 14.2 14.39 14.06 14.83 13.86 14.1 14.12
13.75 14.75 14.38 13.63 14.3 13.83 14.19 13.64 12.93 13.71 12.85 13.5
13.05 13.39 13.3 13.87 14.02 13.73 13.58 13.68 13.76 13.51 13.48 13.28
13.07 14.22 13.56 13.41 13.88 14.21 13.9 13.94 13.82 13.77 13.74 13.29
13.72 12.37 12.33 12.64 13.67 12.17 13.11 13.34 12.21 12.29 13.49 12.99
11.96 11.66 13.03 11.84 12.7 12. 12.72 12.08 12.67 12.16 11.65 11.64
12.69 11.62 12.47 11.81 12.6 12.34 11.82 12.51 12.42 12.25 12.22 11.61
11.46 12.52 11.76 11.41 11.03 12.77 11.45 11.56 11.87 12.07 12.43 11.79
12.04 12.86 12.88 12.81 12.53 12.84 13.36 13.52 13.62 12.87 13.32 13.08
12.79 13.23 12.58 13.17 13.84 12.45 14.34 12.36 13.69 12.96 13.78 13.45
12.82 13.4 12.2 14.16 13.27 14.13]
```

Unique values in 'Malicacid' column:

```
[1.71 1.78 2.36 1.95 2.59 1.76 1.87 2.15 1.64 1.35 2.16 1.48 1.73 1.81
1.92 1.57 1.59 3.1 1.63 3.8 1.86 1.6 2.05 1.77 1.72 1.9 1.68 1.5
1.66 1.83 1.53 1.8 1.65 3.99 3.84 1.89 3.98 4.04 3.59 2.02 1.75 1.67
1.7 1.97 1.43 0.94 1.1 1.36 1.25 1.13 1.45 1.21 1.01 1.17 1.19 1.61
1.51 1.09 1.88 0.9 2.89 0.99 3.87 0.92 3.86 0.89 0.98 2.06 1.33 2.83
1.99 1.52 2.12 1.41 1.07 3.17 2.08 1.34 2.45 2.55 1.29 3.74 2.43 2.68
0.74 1.39 1.47 3.43 2.4 4.43 5.8 4.31 2.13 4.3 2.99 2.31 3.55 1.24
2.46 4.72 5.51 2.96 2.81 2.56 4.95 3.88 3.57 5.04 4.61 3.24 3.9 3.12
2.67 3.3 5.19 4.12 3.03 3.83 3.26 3.27 3.45 2.76 4.36 3.7 3.37 2.58
4.6 2.39 2.51 5.65 3.91 4.28 4.1 ]
```

Unique values in 'Ash' column:

```
[2.43 2.14 2.67 2.5 2.87 2.45 2.61 2.17 2.27 2.3 2.32 2.41 2.39 2.38
2.7 2.72 2.62 2.48 2.56 2.28 2.65 2.36 2.52 3.22 2.8 2.21 2.84 2.55
2.1 2.51 2.31 2.12 2.59 2.29 2.44 2.4 2.04 2.6 2.42 2.68 2.25 2.46
1.36 2.02 1.92 2.16 2.53 1.7 1.75 2.24 1.71 2.23 1.95 2. 2.2 2.58
2.26 2.22 2.74 1.98 1.9 1.88 1.94 1.82 2.92 1.99 2.19 3.23 2.73 2.13
2.78 2.54 2.64 2.35 2.15 2.75 2.69 2.86 2.37]
```

Unique values in '0D280\_0D315\_of\_diluted\_wines' column:

```
[3.92 3.4 3.17 3.45 2.93 2.85 3.58 3.55 2.82 2.9 2.73 3. 2.88 2.65
2.57 3.36 3.71 3.52 4. 3.63 3.82 3.2 3.22 2.77 3.59 2.71 2.87 3.47
2.78 2.51 2.69 3.53 3.38 3.56 3.35 3.33 3.44 2.75 3.1 2.91 3.37 3.26
3.03 3.31 2.84 1.82 1.67 1.59 2.46 2.23 2.3 3.18 3.48 1.93 3.07 3.16
3.5 3.13 2.14 2.48 2.52 2.31 3.12 3.14 2.72 2.01 3.08 2.26 3.21 2.27
2.06 3.3 2.96 2.63 2.74 2.83 2.44 3.57 2.42 3.02 2.81 2.5 3.19 2.12
3.05 3.39 3.69 3.64 3.28 1.29 1.42 1.36 1.51 1.58 1.27 1.69 2.15 2.47
2.05 2. 1.68 1.33 1.86 1.62 1.3 1.47 1.55 1.48 1.64 1.73 1.96 1.78
2.11 1.75 1.56 1.8 1.92 1.83 1.63 1.71 1.74 1.6 ]
```

Unique values in 'Proline' column:

```
[1065 1050 1185 1480 735 1450 1290 1295 1045 1510 1280 1320 1150 1547
1310 1130 1680 845 780 770 1035 1015 830 1195 1285 915 1515 990
1235 1095 920 880 1105 1020 760 795 680 885 1080 985 1060 1260
1265 1190 1375 1120 970 1270 520 450 630 420 355 678 502 510
750 718 870 410 472 886 428 392 500 463 278 714 515 495
562 625 480 290 345 937 660 406 710 438 415 672 315 488
312 325 607 434 385 407 372 564 465 365 380 378 352 466
342 580 530 560 600 650 695 720 590 550 855 425 675 640
725 620 570 615 685 470 740 835 840]
```

```
1 logistic_df
2 categorical = [var for var in logistic_df.columns if logistic_df[var].dtype=='O']
3 print('There are {} categorical variables\n'.format(len(categorical)))
4 print('The categorical variables are :', categorical)
```

There are 0 categorical variables

The categorical variables are : []

## ✓ Cleaning

Nothing to clean. There are no object columns, missing values, and duplicated rows.

## ✓ Logistic Regression

## ✓ Declare Feature Vector and Target Variable

```
1 X = logistic_df.drop('class', axis = 1)
2 y = logistic_df['class']
```

## ✓ Split data into separate training and test set

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
1 X_train.shape, X_test.shape
((142, 13), (36, 13))
```

## ✓ Feature Engineering

```
1 X_train.dtypes

Alcohol                float64
Malicacid              float64
Ash                   float64
Alcalinity_of_ash      float64
Magnesium              int64
Total_phenols          float64
Flavanoids             float64
Nonflavanoid_phenols   float64
Proanthocyanins        float64
Color_intensity        float64
Hue                   float64
0D280_0D315_of_diluted_wines float64
Proline               int64
dtype: object

1 #Display Categorical Variables
2 categorical = [col for col in X_train.columns if X_train[col].dtypes == '0']
3 categorical

[]

1 #Display Numerical Variables
2 numerical = [col for col in X_train.columns if X_train[col].dtypes != '0']
3 numerical

['Alcohol',
 'Malicacid',
 'Ash',
 'Alcalinity_of_ash',
 'Magnesium',
 'Total_phenols',
 'Flavanoids',
 'Nonflavanoid_phenols',
 'Proanthocyanins',
 'Color_intensity',
 'Hue',
 '0D280_0D315_of_diluted_wines',
 'Proline']
```

## ✓ Engineering Missing Values in Numerical Variables

```
1 #Check Missing Values in numerical variables in X_train
2 X_train[numerical].isnull().sum()

Alcohol                0
Malicacid              0
Ash                   0
Alcalinity_of_ash      0
Magnesium              0
Total_phenols          0
Flavanoids             0
Nonflavanoid_phenols   0
Proanthocyanins        0
```

```

Color_intensity      0
Hue                  0
0D280_0D315_of_diluted_wines 0
Proline              0
dtype: int64

```

```

1 #Check Missing Values in numerical variables in X_test
2 X_test[numerical].isnull().sum()

```

```

Alcohol      0
Malicacid    0
Ash          0
Alcalinity_of_ash 0
Magnesium    0
Total_phenols 0
Flavanoids   0
Nonflavanoid_phenols 0
Proanthocyanins 0
Color_intensity 0
Hue          0
0D280_0D315_of_diluted_wines 0
Proline      0
dtype: int64

```

```

1 #Print Percentage of missing values in the numerical variables in training set
2 for col in numerical:
3     if X_train[col].isnull().mean()>0:
4         print(col, round(X_train[col].isnull().mean(), 4))

```

No missing data

## Feature Scaling

```
1 X_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	F
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000
mean	12.984859	2.372606	2.366901	19.554930	100.063380	2.258662	0.497445
std	0.807175	1.115360	0.269684	3.442549	14.249158	0.611691	0.502555
min	11.030000	0.740000	1.360000	10.600000	70.000000	1.100000	0.000000
25%	12.347500	1.602500	2.222500	17.250000	89.000000	1.705000	0.000000
50%	13.040000	1.895000	2.360000	19.500000	98.000000	2.210000	0.000000
75%	13.637500	3.222500	2.560000	21.500000	106.750000	2.735000	0.000000
max	14.750000	5.650000	3.220000	30.000000	162.000000	3.880000	1.000000

```
1 cols = X_train.columns
```

```

1 from sklearn.preprocessing import MinMaxScaler
2 scaler = MinMaxScaler()
3 X_train = scaler.fit_transform(X_train)
4 X_test = scaler.transform(X_test)

```

```
1 X_train = pd.DataFrame(X_train, columns = [cols])
```

```
1 X_test = pd.DataFrame(X_test, columns = [cols])
```

```
1 X_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	F
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	
mean	0.525500	0.332506	0.541345	0.461594	0.326776	0.416785	
std	0.216983	0.227161	0.144991	0.177451	0.154882	0.220033	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.354167	0.175662	0.463710	0.342784	0.206522	0.217626	
50%	0.540323	0.235234	0.537634	0.458763	0.304348	0.399281	
75%	0.700941	0.505601	0.645161	0.561856	0.399457	0.588129	
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

## Model Training

```

1 #Train a Logistic regression model on the training set
2 from sklearn.linear_model import LogisticRegression
3
4 #Instantiate the model
5 logreg = LogisticRegression(solver = 'liblinear', random_state = 0)
6
7 #fit the model
8 logreg.fit(X_train, y_train)

```

```

LogisticRegression
LogisticRegression(random_state=0, solver='liblinear')

```

## Predict Results

```

1 y_pred_test = logreg.predict(X_test)
2 y_pred_test

array([1, 3, 2, 1, 2, 2, 1, 3, 2, 2, 3, 3, 1, 2, 3, 2, 1, 1, 3, 1, 2, 1,
       1, 2, 2, 2, 2, 2, 3, 1, 1, 2, 1, 1])

```

predict\_proba method gives the probabilities for the class, or origin of wine, which in this case is 1, 2, and 3, in array form.

```

1 #Probability of getting output as 1
2 logreg.predict_proba(X_test)[: ,0]

array([0.83411707, 0.08169456, 0.33794154, 0.79491944, 0.2372458 ,
       0.237406 , 0.87096795, 0.03861458, 0.16257044, 0.05946295,
       0.1482802 , 0.03716879, 0.92972498, 0.47744117, 0.0838979 ,
       0.13261824, 0.78069736, 0.95517064, 0.0825623 , 0.84310019,
       0.47020366, 0.67960296, 0.47193266, 0.24693731, 0.08366256,
       0.16146061, 0.21250699, 0.05442331, 0.07420845, 0.07434986,
       0.83480123, 0.85379414, 0.08617887, 0.8224449 , 0.87731904,
       0.6611851 ])

```

```

1 #Probability of getting output as 2
2 logreg.predict_proba(X_test)[: ,1]

array([0.1240818 , 0.06966025, 0.65228402, 0.1607158 , 0.61192177,
       0.75455152, 0.08070117, 0.14839743, 0.77629276, 0.77373406,
       0.13986198, 0.07780088, 0.03752654, 0.51475621, 0.07579536,
       0.85273997, 0.16126423, 0.02246023, 0.44592902, 0.13936754,
       0.52143812, 0.24521823, 0.43573301, 0.71879674, 0.58483752,
       0.76777135, 0.73705714, 0.8411666 , 0.73358901, 0.06019109,
       0.12898216, 0.11348826, 0.61642905, 0.06319566, 0.08703462,
       0.31648131])

```

```

1 #Probability of getting output as 3
2 logreg.predict_proba(X_test)[: ,2]

array([0.04180113, 0.84864519, 0.00977444, 0.04436476, 0.15083243,
       0.00804248, 0.04833088, 0.81298799, 0.0611368 , 0.16680299,
       0.71185782, 0.88503033, 0.03274848, 0.00780262, 0.84030674,

```

```
0.01464178, 0.0580384 , 0.02236913, 0.47150868, 0.01753227,
0.00835821, 0.0751788 , 0.09233433, 0.03426595, 0.33149992,
0.07076804, 0.05043587, 0.10441009, 0.19220254, 0.86545905,
0.03621662, 0.0327176 , 0.29739208, 0.11435944, 0.03564634,
0.02233359])
```

## ✓ Check Accuracy Score

```
1 from sklearn.metrics import accuracy_score
2 print('Model accuracy score: {0:0.4f}'. format(accuracy_score(y_test, y_pred_test)))
```

```
Model accuracy score: 0.9722
```

## ✓ Compare the train-set and test-set accuracy

```
1 y_pred_train = logreg.predict(X_train)
2 y_pred_train

array([3, 2, 2, 3, 1, 1, 2, 2, 2, 1, 3, 2, 3, 1, 3, 3, 1, 3, 1, 2, 3, 3,
       2, 3, 3, 1, 3, 3, 2, 3, 3, 2, 1, 2, 2, 2, 1, 1, 2, 2, 3, 3, 2, 2,
       2, 3, 3, 1, 3, 2, 2, 2, 2, 2, 1, 1, 2, 1, 3, 1, 3, 1, 1, 2, 1, 2,
       2, 1, 3, 2, 1, 2, 2, 2, 3, 1, 3, 3, 1, 1, 2, 3, 1, 1, 2, 2, 1, 1,
       1, 3, 2, 1, 2, 3, 1, 2, 3, 3, 1, 1, 3, 1, 3, 2, 1, 1, 2, 1, 3, 2,
       3, 1, 3, 3, 3, 1, 2, 2, 2, 3, 3, 2, 2, 2, 2, 3, 3, 1, 1, 3, 2,
       2, 2, 1, 1, 1, 2, 2, 2, 1, 3])
```

```
1 print('Training-set accuracy score: {0:0.4f}'. format(accuracy_score(y_train, y_pred_train)))
```

```
Training-set accuracy score: 0.9789
```

## ✓ Check for overfitting and underfitting

```
1 #Print the scores on training and test set
2 print('Training set score: {:.4f}'. format(logreg.score(X_train, y_train)))
3 print('Test set score: {:.4f}'. format(logreg.score(X_test, y_test)))
```

```
Training set score: 0.9789
Test set score: 0.9722
```

As the training set score and test set score are quite comparable to each other, there is no question of overfitting.

## ✓ Use C = 100

```
1 #Fit the Logistic Regression model with C = 100
2 #Instantiate the model
3 logreg100 = LogisticRegression(C = 100, solver = 'liblinear', random_state = 0)
4 #Fit the model
5 logreg100.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression(C=100, random_state=0, solver='liblinear')
```

```
1 #Print the scores on training and test set
2 print('Training set score: {:.4f}'.format(logreg100.score(X_train, y_train)))
3 print('Test set score: {:.4f}'.format(logreg100.score(X_test, y_test)))
```

```
Training set score: 1.0000
Test set score: 1.0000
```

We can see that C = 100 results in higher accuracy, meaning that this model performed the best. But let us be careful and use other means, in case that this is a result of overfitting.



```

1 #Fit the Logistic Regression model with C = 0.01
2 #Instantiate the model
3 logreg001 = LogisticRegression(C = 0.01, solver = 'liblinear', random_state = 0)
4 #Fit the model
5 logreg001.fit(X_train, y_train)

```

```

▼                LogisticRegression
LogisticRegression(C=0.01, random_state=0, solver='liblinear')

```

```

1 #Print the scores on training and test set
2 print('Training set score: {:.4f}'.format(logreg001.score(X_train, y_train)))
3 print('Test set score: {:.4f}'.format(logreg001.score(X_test, y_test)))

```

```

Training set score: 0.8944
Test set score: 0.8889

```

Setting a different value for C other than 100 actually prevents the model from overfitting, as we can see from when C is equals to 1 and 0.01.

## ▼ Compare model accuracy with null accuracy

```

1 #Check class distribution in test set
2 y_test.value_counts()

```

```

class
2    16
1    14
3     6
Name: count, dtype: int64

```

We can see that the occurrences of the most frequent class is 16. So, we can calculate the null accuracy by dividing 16 by the total number of occurrences

```

1 #Check Null Accuracy Score
2 null_accuracy = (16/(16+14+6))
3 print('Null Accuracy Score: {:.4f}'.format(null_accuracy))

```

```

Null Accuracy Score: 0.4444

```

As the scores for the training and test when C is equal to one are 0.9789 and 0.9722, respectively, a null accuracy score of 0.4444—which is significantly lesser than both score—tells us that the model is performing exceptionally well and is learning meaningful patterns from the data.

## ▼ Confusion Matrix

```

1 #Print the Confusion Matrix and slice it into four pieces
2 from sklearn.metrics import confusion_matrix
3 cm = confusion_matrix(y_test, y_pred_test)
4 print('Confusion matrix\n\n', cm)
5 print("\nTrue Positives(TP) = ", cm[0,0])
6 print("\nTrue Negatives(TN) = ", cm[1,1])
7 print("\nFalse Positives(FP) = ", cm[0,1])
8 print("\nFalse Negatives(FN) = ", cm[1,0])

```

```

Confusion matrix

```

```

[[14  0  0]
 [ 0 15  1]
 [ 0  0  6]]

```

```

True Positives(TP) = 14

```

```

True Negatives(TN) = 15

```

```

False Positives(FP) = 0

```

```

False Negatives(FN) = 0

```

The confusion matrix shows 19 correct predictions, and 0 incorrect predictions!!

```

1 #Visualize confusion matrix with seaborn heatmap
2 cm_matrix = pd.DataFrame(data = cm, columns=['Predicted Class 1', 'Predicted Class 2', 'Predicted Class 3'],
3                             index=['Actual Class 1', 'Actual Class 2', 'Actual Class 3'])
4 sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')

```

